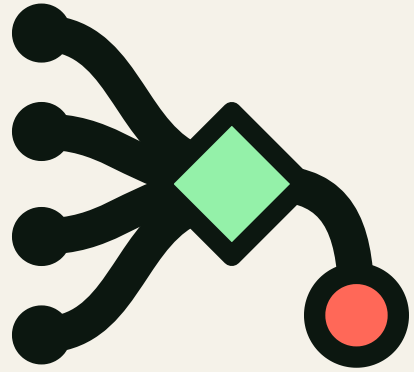
OPEN-SOURCE EXPERIMENT

Jeviphy



Source-aware code search from reviewed memory

Graphify remains the graph engine

THE PROBLEM

Broad questions can rank unrelated nodes before the right source



Question › Reviewed source paths › Graphify traversal › Source-cited result

Memory narrows the search. The agent still verifies the cited source.

LOCAL WORKFLOW

Five bounded steps, with an exact fallback



1

Confirm graph and repository commit match

2

Read selected, reviewed observations

3

Validate source paths in the repository

4

Seed traversal with the best three nodes

5

Fall back when evidence is absent or invalid

The measured `fast-source` path runs offline. Paid Jev calls stay optional and off by default.

FROZEN CLICK COHORT · 40 QUESTIONS · 5 TIMING PASSES

Jeviphy improved retrieval and latency in the measured experiment



198 of 200 paired source-backed runs were faster. Control answers stayed byte-identical.

Same AST graph, source commit, machine, order policy and question set

FAILED EXPERIMENT

The paid reranker did not clear its advancement gate



4 / 19

Top-five hits stayed unchanged on identical bounded candidate pools

Added API latency

649.74 ms

p50

1,083.09 ms

p95

Jev remains optional. The positive result came from local source-cited retrieval.

EVIDENCE BOUNDARY

What the current experiment supports



Supported

- Promising results for synthetic source-cited questions
- Better retrieval and latency on one frozen Click cohort
- Offline operation with the original Graphify fallback
- A reproducible 60-second demo with no API key

Still open

- Real-user task outcomes
- Relevance gains on real project memory
- Results across more repositories and question types
- Interoperability with Claude Mem and upstream Graphify

REVIEW REQUEST

Three questions for maintainers and testers



Graphify maintainers

Should source-seeded traversal live inside Graphify or remain an external adapter?

Claude Mem maintainers

Is an optional, reviewed `source_paths` field a useful interoperability boundary?

Independent testers

Can the frozen workflow reproduce on another repository without weakening fallback or privacy?

Start with the 60-second offline demo, then inspect the independent comparison.